

Silent Witness — Project Intro Script

(~4-5 minute spoken script, plain language, no jargon-first explanations)

Elevator Pitch (10-15 seconds)

"Imagine five people watch the same car accident. Later, each of them tells a slightly different story — different colors, different times, different directions. Right now, a journalist or a lawyer has to manually read all five statements and try to spot the differences by hand. **Silent Witness is an AI system that reads multiple eyewitness statements of the same event, and automatically shows you exactly where they agree and where they contradict each other — with proof, on a timeline and a map. It never decides who's lying. It just shows you the facts, side by side.**"

1. The Problem (30 seconds)

"When something happens — a robbery, an accident, a fire — you rarely get just one account of it. You get three, five, sometimes ten. Every witness saw it from a different angle, remembers different details, and describes it in their own words.

Today, someone — a journalist building a story, or a legal reviewer preparing a case — has to read every single statement, and manually cross-check them. Was the car red or blue? Did it happen at 3pm or 3:15? Was there one person or two? This is done today on whiteboards and spreadsheets. It's slow, it doesn't scale, and it's easy to miss a contradiction buried on page four of statement seven."

2. The Solution — What Silent Witness Does (30 seconds)

"Silent Witness automates that cross-checking step. You feed it multiple witness statements about the same incident. It reads them, pulls out the *facts* each person is claiming — who, what, where, when — and lines those facts up against each other.

Then it shows you two things clearly: where witnesses **agree** with each other, which strengthens a claim, and where they **contradict** each other, which is exactly what a reviewer needs to dig into. Every single flag — agreement or contradiction — links straight back to the exact sentence it came from, in the exact statement it came from. No black box. No 'trust the AI.' You can always click through and read it yourself."

3. The One Rule We Never Break (15 seconds)

"There's one thing this system deliberately does **not** do: it never tells you who's lying. It doesn't score witnesses as reliable or unreliable. It only ever says 'these two facts don't match' — the human reviewer

decides what that means. That's a deliberate ethical boundary, not a limitation. An 'AI lie detector' is both unreliable and dangerous to build — a fact-comparison tool is genuinely useful and honest about what it can and can't tell you."

4. How the ML Model Actually Works, in Plain Language (60-90 seconds)

"So how does the AI actually do this? Think of it as a five-step reading process, the same way a very careful human reviewer would work — just automated and consistent every time.

Step one — Reading and tagging. For every statement, the system reads the text and picks out the important pieces: people, vehicles, objects, places, and times. It's like highlighting a document with different colored pens — this pen for 'who,' this pen for 'when,' this pen for 'where.'

Step two — Understanding what happened. Beyond just tagging words, it figures out the actual *events* — who did what, to what, and when. Not just 'red car' and '3pm' as separate facts, but 'the red car turned left at 3pm.'

Step three — Matching up the same things across witnesses. This is the clever part. Witness A might say 'the car.' Witness B might say 'the red sedan.' Witness C might just say 'it.' The system has to figure out — are these all talking about the *same* car? It uses AI language understanding to link these mentions together, the same way you'd mentally connect 'the car' and 'it' while reading a story.

Step four — Comparing the facts. Once it knows which claims are talking about the same thing, it compares them. If the facts match — that's a corroboration. If they conflict — wrong color, wrong direction, wrong time, or even 'this witness says it never happened at all' — that's flagged as a contradiction. Some of this comparison is simple rule-checking, like comparing two colors directly. Some of it needs a smarter AI model, because people describe the same thing in very different words, and the system has to understand that 'sped away' and 'drove off quickly' mean the same thing.

Step five — Showing it visually. Finally, everything gets laid out on a timeline — so you see the whole incident unfold in order — and on a map — so you see where every claim happened, geographically. Contradictions and agreements are visually marked right there. Click anything, and it jumps you straight back to the original sentence it came from."

5. The System Flow, End to End (30-45 seconds)

"Here's what that looks like in practice, from a reviewer's point of view: they open the tool, paste in or upload the witness statements. Behind the scenes, the system processes each one — usually in just a few seconds — extracting facts, matching them up, and comparing them.

The reviewer then sees a dashboard: a timeline of the whole incident, a map of where things happened, and a clear list of contradictions and agreements. They can click any flagged item and immediately see the exact original wording from both witnesses side by side. At the end, they export a clean report — ready to go into an article, a case file, or an insurance review."

6. Why This Is Technically Solid (20 seconds, optional — for a technical audience)

"Under the hood, this isn't one big model trying to do everything — it's a pipeline of specialized steps, each one testable and improvable on its own: entity extraction, event understanding, cross-witness matching, and contradiction detection, backed by a synthetic dataset with full ground truth so we can actually measure accuracy, not just eyeball it. The whole thing is built as a scalable, modular system — the AI processing runs separately from the web app, so it can grow independently as the system handles more incidents and more witnesses."

7. Closing Line (10 seconds)

"In short: Silent Witness doesn't replace human judgment — it removes the tedious, error-prone part of the job, so reviewers spend their time on what actually matters: deciding what the truth is, with all the facts laid out clearly in front of them."

Speaker Notes

Sections 1–3 and 7 are the "must keep" core if you're cut for time.

Section 4 is the heart of the technical explanation — slow down here, use hand gestures for the "five steps," and it's the natural place to click through to a live demo screen if you have one.

Section 6 can be skipped entirely for a non-technical audience (journalists, general supervisors) and kept for a technical panel/defense.